

MADHAV S

AI & Data Science Student

github.com/smadhav2024

PROJECT REVIEW REPORT

1. EXECUTIVE SUMMARY

This report provides a comprehensive review of two major enterprise-grade Artificial Intelligence initiatives: the "Real-Time 3D AI Avatar Interface" and "Forge," a self-hosted Small Language Model (SLM) training platform. Both projects were engineered from the ground up to solve complex human-computer interaction and data-processing bottlenecks, delivering highly optimized, production-ready AI solutions.

2. PROJECT I: REAL-TIME 3D AI AVATAR INTERFACE

2.1 Project Statement & Objective

The primary objective was to architect a high-performance, interactive 3D AI Assistant. The legacy codebase suffered from severe latency, jittery animations, and rendering crashes that degraded the user experience. The goal was to build a real-time, low-latency interface featuring fluid 3D body language, mathematically accurate lip-syncing, and instant large language model (LLM) streaming. The final product operates as an emotionally intelligent Virtual Assistant capable of discussing domain-specific knowledge and answering complex queries with natural human pacing.

2.2 Target Beneficiaries

- **End Users & Customers:** Gain immediate, conversational access to complex information, support, and services without waiting for human representatives.
- **Support & Operations Teams:** Reduces the administrative burden of answering repetitive inquiries, allowing staff to focus on complex, high-value problem-solving.
- **Enterprise Organizations:** Enhances product adoption and user engagement through an intuitive, emotionally intelligent, and highly accessible interface.

2.3 Technical Implementation & Innovations

- **Streaming Audio Engine:** Engineered a custom Web Audio API backend leveraging FastAPI and WebSockets. The system decodes base64 PCM audio chunks on the fly from the

Cartesia AI TTS engine, carefully managing buffer synchronization to eliminate playback latency and stuttering.

- **Advanced Lip-Syncing Architecture:** Replaced legacy phoneme systems with a procedural grapheme-to-viseme engine. Text syllables are dynamically mapped to 32 standard ARKit/Oculus visemes. Morph targets are mathematically interpolated (lerping) frame-by-frame, ensuring flawless, realistic mouth movements.
- **Dynamic Animation & State Management:** Restored over 20 Mixamo skeletal animations (e.g., Talking, Thoughtful Nod, Relieved Sigh). Engineered a "Pendulum" breathing system and state tracker that forces the avatar to return to a natural resting state after gestures, eliminating animation jitter. A Smart Animation Randomizer ensures non-repetitive body language.
- **Context-Aware Expressions:** Integrated real-time text sentiment analysis to dynamically drive facial morph targets. The avatar instantly smiles, frowns, or looks thoughtful based on the semantic context of the generated response.

3. PROJECT II: FORGE - SELF-HOSTED SLM TRAINING PLATFORM

3.1 Project Statement & Objective

Data quality is the most critical bottleneck in AI fine-tuning. The objective of the Forge platform was to solve the "Garbage In, Garbage Out" problem by building a full-stack, localized environment. Forge allows users to upload highly unstructured datasets (CSVs, JSONs, PDFs, Word Docs) and automatically clean, deduplicate, and format them for Low-Rank Adaptation (LoRA) fine-tuning. By completely containerizing the pipeline, the platform ensures total data privacy while enabling the precise fine-tuning of advanced small language models (e.g., configuring unsloth/Qwen3-1.7B variants) optimized for any specialized organizational domain.

3.2 Target Beneficiaries

- **AI Engineers & Data Scientists:** Drastically reduces the engineering hours spent on manual data cleaning, text chunking, and dataset formatting.
- **Domain Experts:** Empowers non-technical subject matter experts to easily fine-tune specialized language models on proprietary company data without writing code.
- **Data Privacy & Security Teams:** Ensures absolute data sovereignty by running the entire LLM pipeline locally; sensitive enterprise data never leaves the internal network.

3.3 Technical Implementation & Innovations

- **The 6-Layer Data Pipeline:** Developed a deterministic FastAPI backend that automatically ingests unstructured files. It features automated schema detection, MinHash LSH fuzzy deduplication, perplexity screening to drop corrupted text, and strict regex-based PII scrubbing to redact sensitive information.
- **Synthetic Q&A Generation:** For unstructured prose, the pipeline chunks the text and utilizes a localized 1B SLM bound by strict Pydantic JSON schemas. This mechanism synthetically generates highly structured Instruction-Response pairs without generating hallucinations or conversational filler.
- **Interactive Review Dashboard:** Built a Next.js/React frontend providing users with a "safety net" before model training. Users can dynamically adjust cleaning aggressiveness, manually edit generated Q&A pairs, and restore flagged items from a soft-delete Quarantine Bin.
- **Containerized Training & Evaluation:** Orchestrated HuggingFace SFTTrainer with dynamic LoRA configuration inside Docker. Integrated a side-by-side A/B testing interface that allows users to instantly compare prompt outputs from the base model against the fine-tuned model.

4. KEY LEARNINGS & IMPACT

Throughout the development of these architectures, several critical insights were gained regarding system optimization and user experience. Managing the browser's main thread became a focal point in the 3D Avatar project; deferring heavy background asset loading (via FBXLoader) and clamping render-loop deltas proved essential to preventing frame drops and visual stuttering. Similarly, the Forge platform underscored the necessity of robust data sanitization before fine-tuning, demonstrating that automated pipelines must include human-in-the-loop oversight (like the "Quarantine Bin") to prevent catastrophic data loss. Ultimately, both projects reinforced the importance of balancing computationally heavy backend AI tasks with seamlessly optimized frontend architectures to deliver responsive, enterprise-ready platforms.

5. CONCLUSION

Both the Real-Time 3D Avatar and the Forge SLM Platform represent major technical achievements in AI interface design and infrastructure. The 3D Avatar significantly lowers the barrier to human-AI interaction through sophisticated psychological design, natural movement, and low-latency audio streaming. Concurrently, Forge solves the critical engineering hurdle of automated data sanitization for customized, localized model training. Deployed together, these architectures provide a comprehensive, secure, and highly engaging ecosystem for leveraging domain-specific Artificial Intelligence in a modern enterprise environment.